

# Java-based Simulation of Load Balancing Policies

...

Vaibhav Krishnakumar

June 2018

*Supervisor: Giuliano Casale*

*Second Marker: Alex Carver*

# Distributed Systems

Scalability + Availability

Response Time / Latency

Optimal Weights

Load Balancing



# Model

## ‘Supermarket’ Model

Heterogeneous servers

Multiple job classes

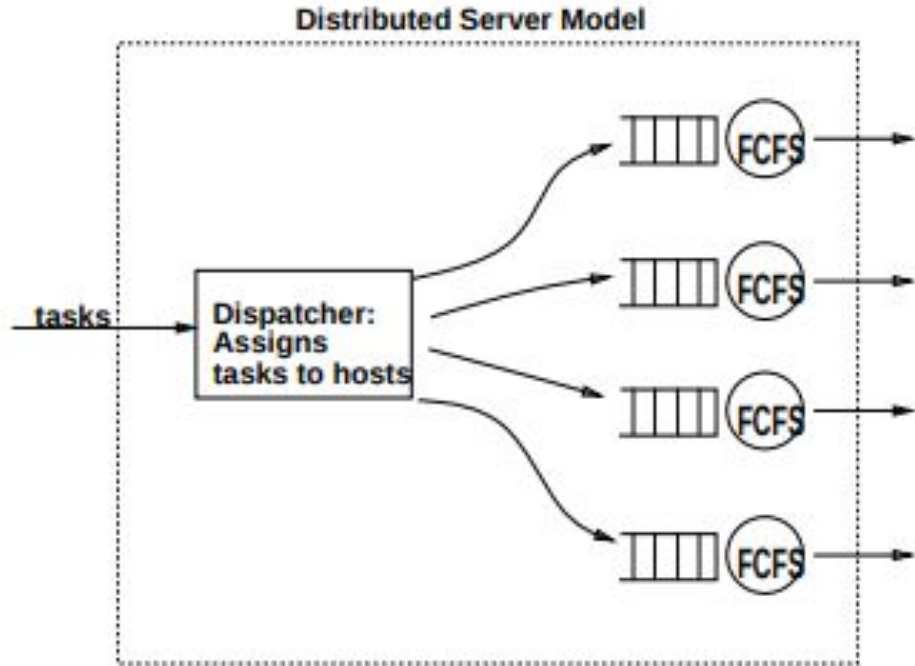
Non preemptive tasks

One-shot dispatch

No redundant tasks

FCFS

$\infty$  capacity



# Java Modelling Tools (JMT)

Politecnico di Milano, Imperial College

Capacity Planning, Queueing Network Simulations

Changes to codebase

UI

XML Model

Routing Logic

50,000 users; 150,000 lines of code



**Imperial College  
London**

# Load Balancing Policies

Random

Round Robin


Weighted Round Robin

Shortest Queue (JSQ)

Power of k

Least Work Left

Size-Interval Task Assignment



# Power of $k$

Performance optimisation to JSQ

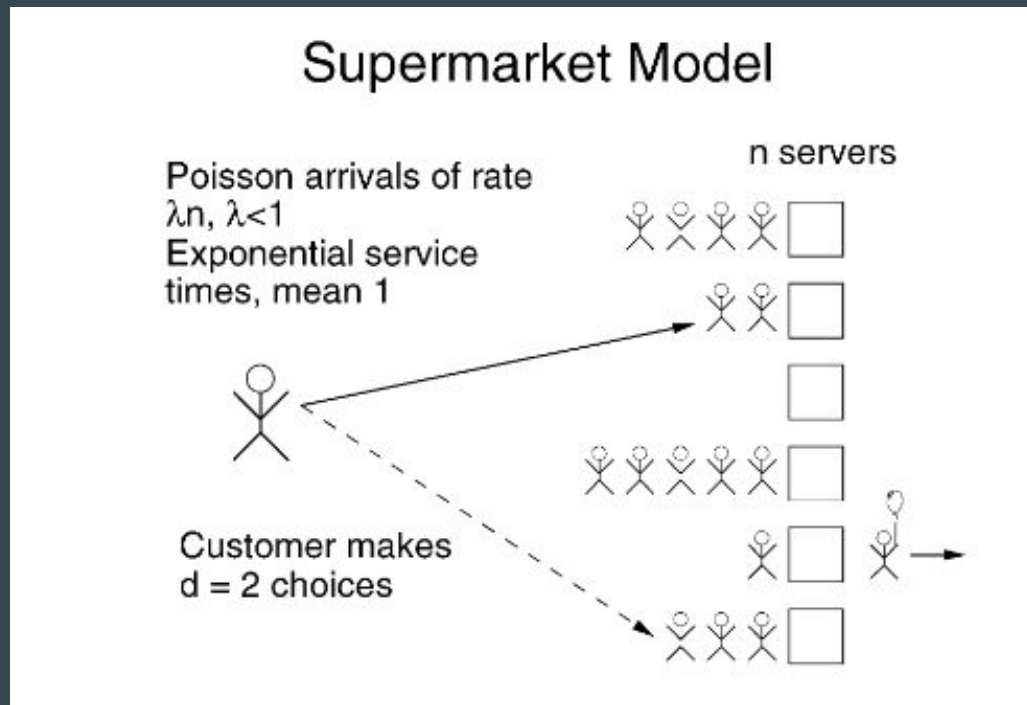
Query  $k < n$  servers on  
each assignment

Send job to shortest observed queue

Trade-off optimality for performance

$k = 1$ : Random

$k = n$ : JSQ



# Power of $k$ with Memory


Add local memory to Power of  $k$

Query  $k < n$  servers on each assignment as before

Store the results in local 'memory'

Send task to best node in memory

*It is possible to send the job to a node not queried in that round*



# Markov Chain Verification

Analytically solve a small, closed model with Power of k algorithms


$N = 4$  for Power of k,  $N = 2$  for 'with Memory'

Solve flow-balance equation and normalise to obtain steady-state distribution

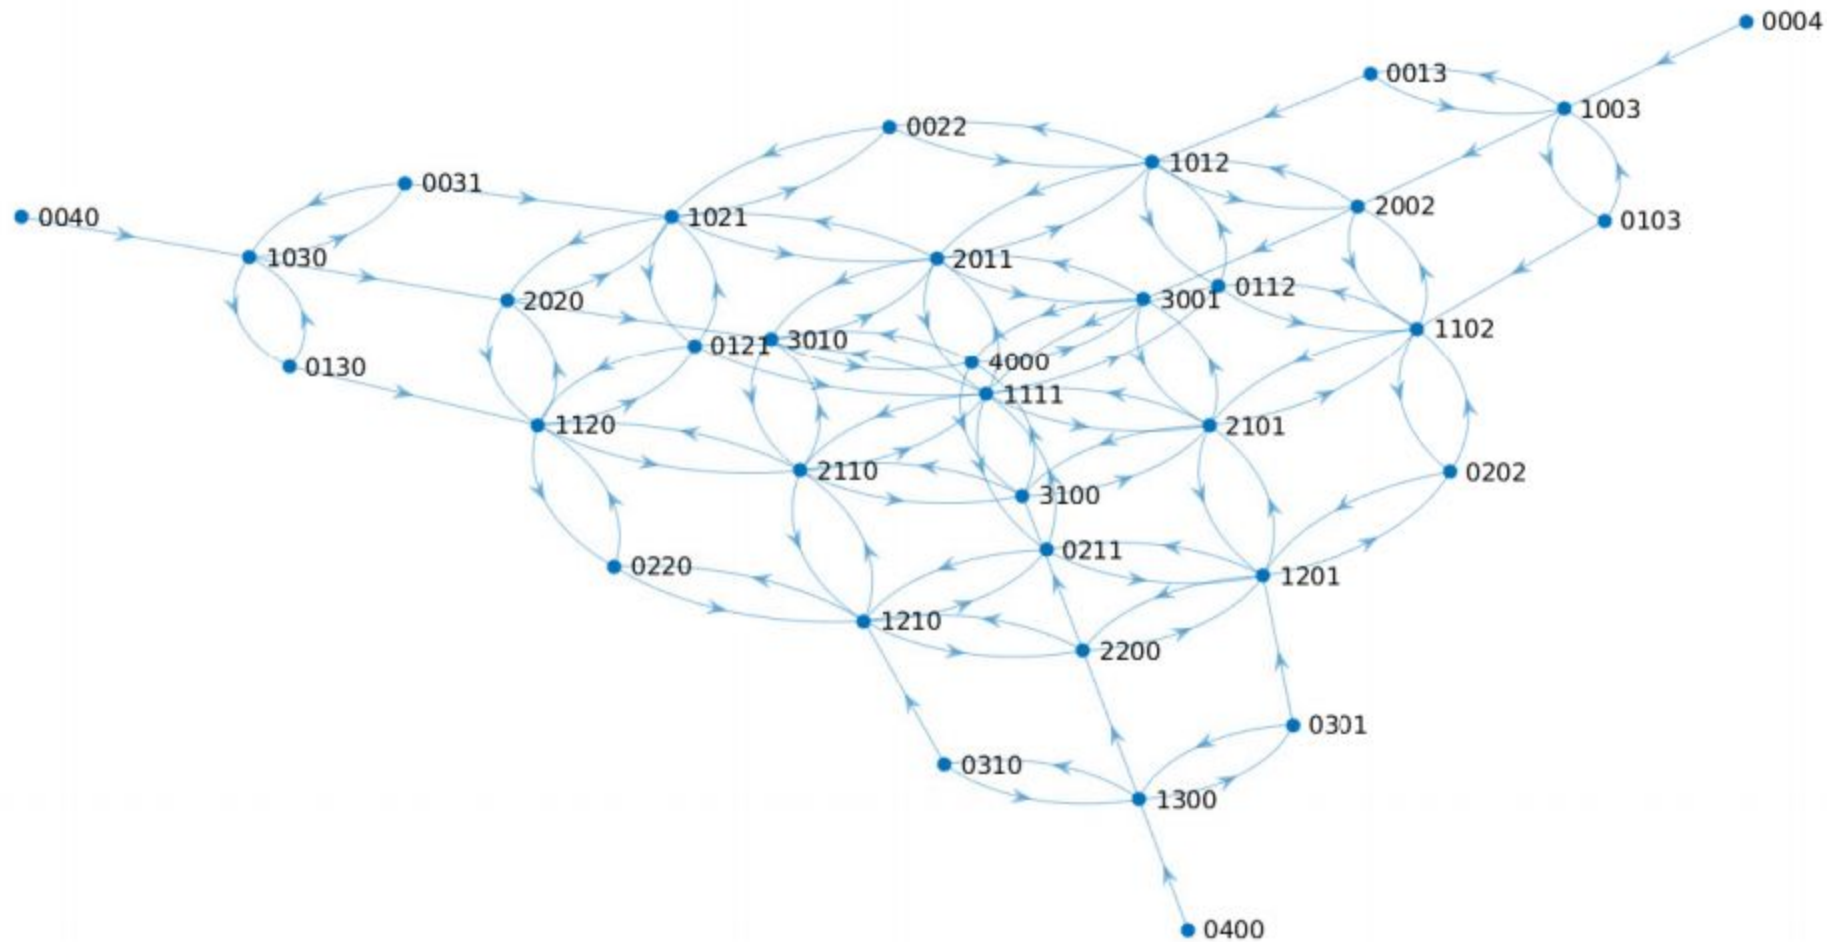
Calculate average queue length, utilisation and response time

Run simulation on JMT and compare

Results match closely, relative error under 1%







# Weighted Round Robin

Cycle through servers, assigning jobs in proportion to weights

Weights are integers representing the number of jobs received in one full execution

Useful to exploit the heterogeneity of machines

Computing optimal weights is difficult in practise

Routing Options	
Destination	Weights
Queue 1	1
Queue 2	2
Queue 3	3

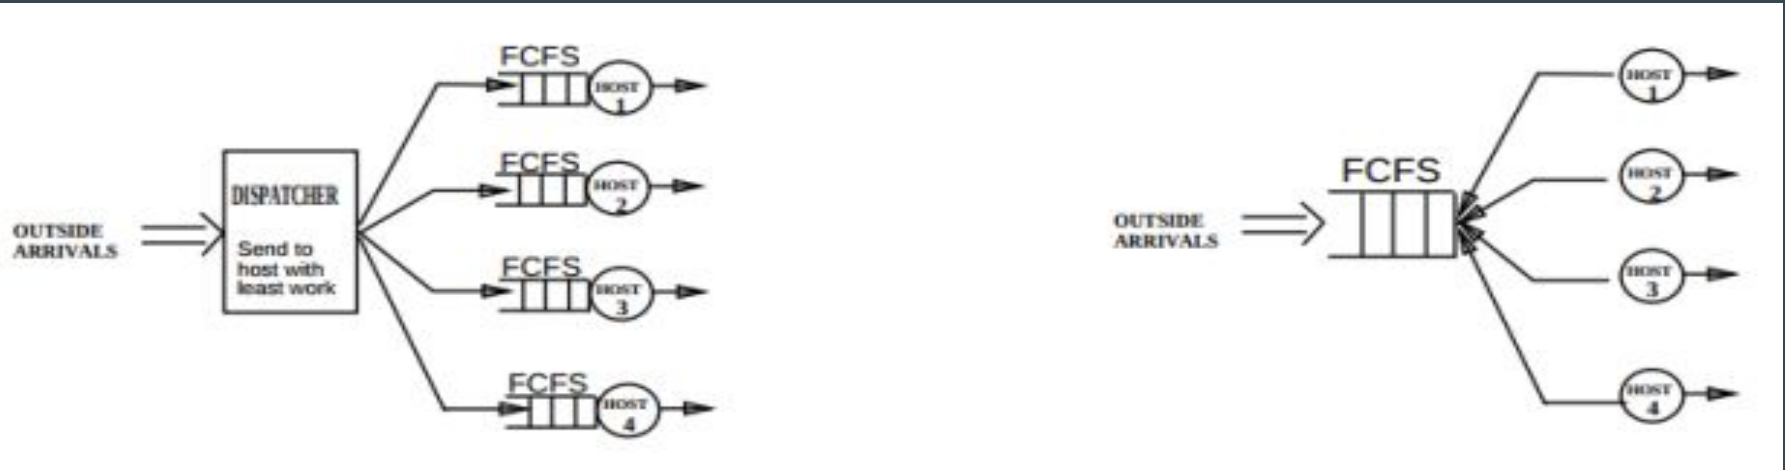
# Least Work Left

Calculate the  $WORK\_LEFT$  metric for each server

Central Queue

Sum of service times of jobs waiting in queue

Residual service time of job in service



# SITA

Partitioning the job size distribution

Allocate a server per partition


Cut-offs vary by objective

SITA-OPT

SITA-E

Bounded Pareto Distribution

File size on the internet




# The impact of adding memory


JSQ is the bounding case of Power of  $k$  strategies

The performance converges as number of choices ( $k$ ) or number of jobs ( $N$ ) increase

Response Time v  $k$



Throughput v  $N$

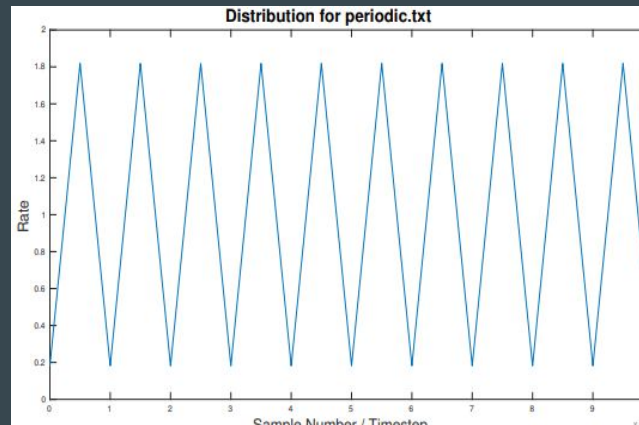
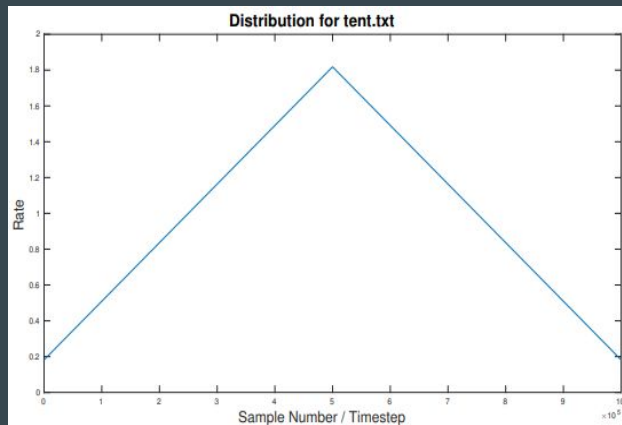
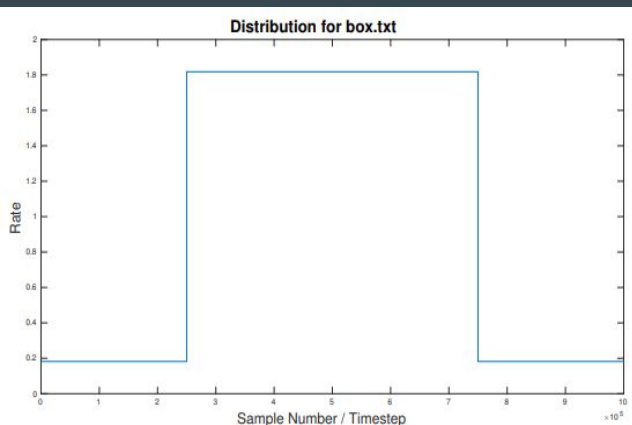


# Resilience to Inter-arrival Times

Vary inter-arrival rates based on distributions below

LWL and Power of  $k$  with Memory reacted best to the changes

Other algorithms perform similarly, further research required to differentiate



# Criteria

Response Time Minimisation

Throughput Maximisation

Utilisation Maximisation

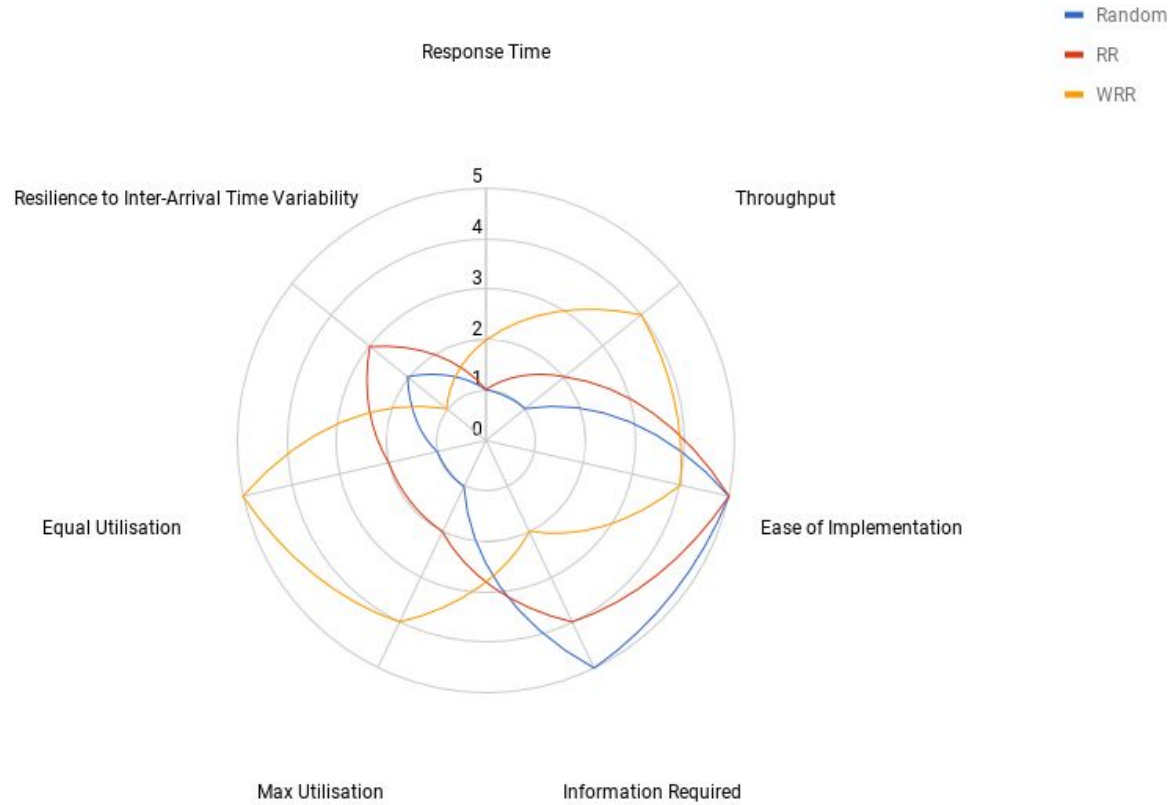
Utilisation Equalisation

Resilience to Inter-arrival Time Variability

Information Required


Ease of Implementation

## Static Policies





## Dynamic Policies



# Evaluation

Power of  $k$  - Markov Chain modelling

Power of  $k$  with Memory

Demonstrate results from original paper

Prof. Anselmi ran further tests, performs as expected

Least Work Left

Further testing needed

Verification is not 'convincing'

# Future Work

Automate Markov Chain generation

- Simple program to create transition matrices

- Larger, more complex models can be analysed with ease

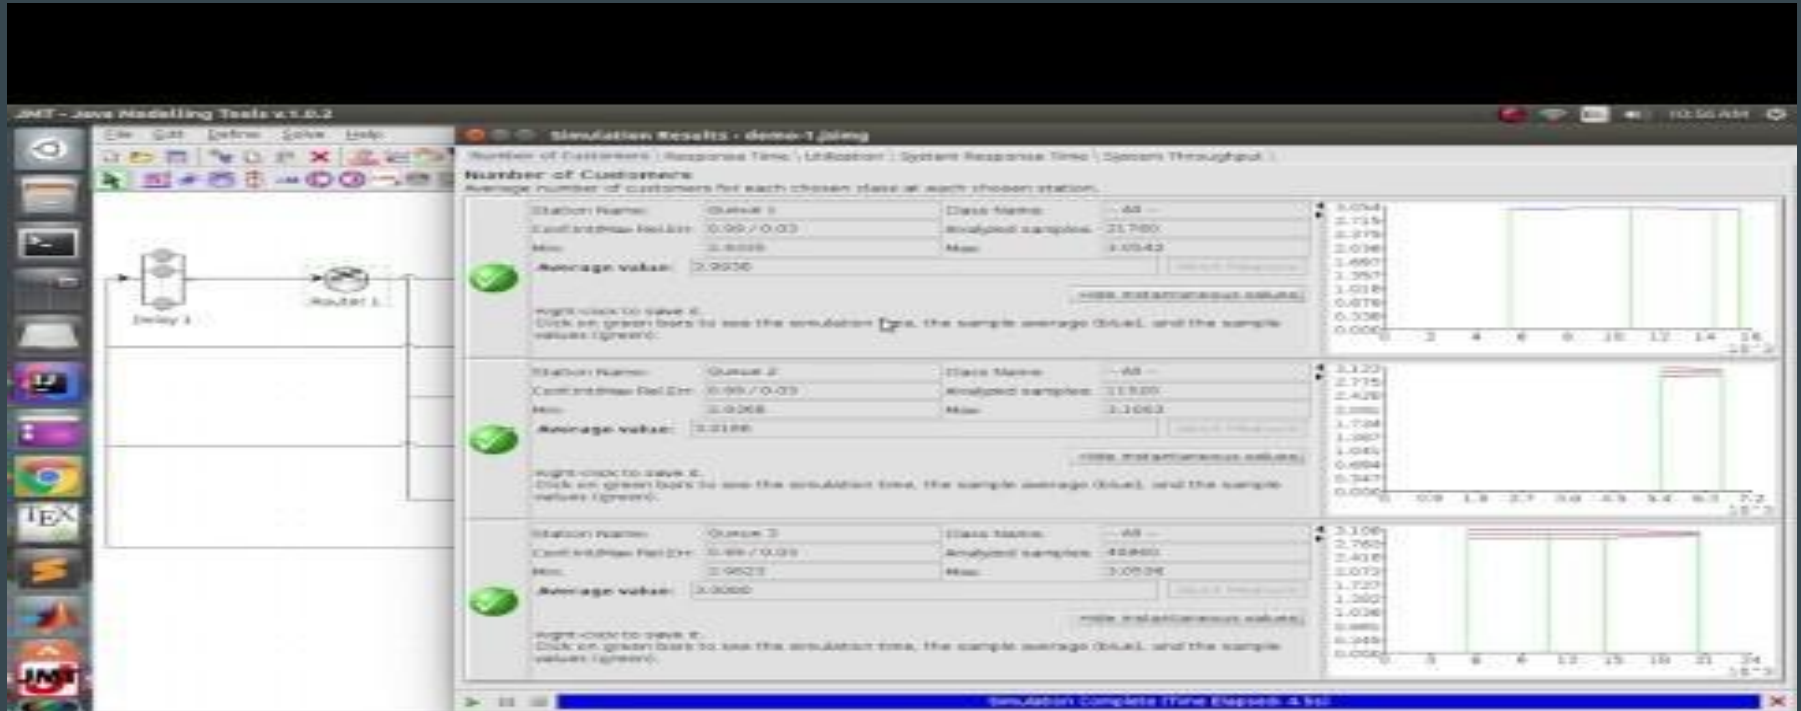
SITA

- Integrate into JMT properly


- Introduce concept of job size

- Collaboration of servers

# Demo - Power of k, Power of k with Memory



# Demo - Weighted Round Robin, Least Work Left



# Demo - SITA

